

# *AN INTRO TO BAYESIAN ML*

**Prof. Chris Jermaine**  
**cmj4@cs.rice.edu**

# Before We Begin

- In A8, implement a ML algorithm to extract “topics” from text
- This is not a ML class... so to be fair
  - Will try to specify algorithm so precisely in the next few lectures
  - That you can implement it without really understanding what is going on
- That said...
  - It would be a shame if this is what happend!
  - So will spend considerable time trying to explain what’s going on
  - And I hope it’s gonna make sense!
- So sit back, enjoy, and hopefully you’ll learn something!

# Modern Machine Learning


- Is a whole lot like applied statistics
- In fact, I'd argue that CS in general is becoming more statistical
  - That is, based on observation and inference
  - ...and a bit less logic-based
- Ask Google, Facebook, LinkedIn how important stats is!

# The Bayesian Approach

- One branch of machine learning utilizes the “Bayesian” approach
- Say our goal was to give E2 to a few students
  - Then use their performance to figure out what the class average will be
  - Want to figure out if E2 is too hard or too easy
- How would a Bayesian do this?

# So How Would a Bayesian Do This?

- First imagine a stochastic “generative process” for the data
  - For the  $i$ th student, we might imagine  $\text{score}_i \sim \text{Normal}(\mu, \sigma^2)$




- Why normal? Models typical “bell shaped” data
- Assume  $\mu, \sigma^2$  are also generated by sampling from RVs, and are unknown

# So How Would a Bayesian Do This?

- First imagine a stochastic “generative process” for the data

— For the  $i$ th student, we might imagine  $\text{score}_i \sim \text{Normal}(\mu, \sigma^2)$

ultimate goal is to  
guess the value of  $\mu$




Means “is  
sampled from” a  
RV with a  
normal dist.

— Why normal? Models typical “bell shaped” data

— Assume  $\mu, \sigma^2$  are also generated by sampling from RVs, and are unknown

# So How Would a Bayesian Do This?

- How is the mean  $\mu$  generated?
  - We imagine  $\mu \sim \text{Normal}(70, 100)$




- Why Normal (70, 100)? Allows for all possible, reasonable exam averages

# So How Would a Bayesian Do This?

- How is the variance (spread)  $\sigma^2$  generated?

— We imagine  $\sigma^2 \sim \text{InverseGamma}(1, 1)$



— Why InverseGamma (1, 1)? Allows a really large range of positive  $\sigma^2$  vals



# Thus, Our Generative Process Is

- Step 1:  $\sigma^2 \sim \text{InverseGamma}(1, 1)$
- Step 2:  $\mu \sim \text{Normal}(75, 100)$
- Step 3: for each  $i$ ,  $\text{score}_i \sim \text{Normal}(\mu, \sigma^2)$

# Why Have a Generative Process?

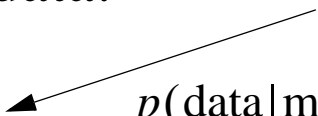
- Given gen process, can measure how likely a given pair is
  - Specifically,  $p(\mu, \sigma^2) = \text{IG}(\sigma^2 \mid 1, 1) \times \text{Normal}(\mu \mid 75, 100)$
  - So given any  $\mu, \sigma^2$  combo, we can say how likely we think it is
- This is our “prior” on mean and variance

Note “times”  
cause the two  
values are  
indep.

# Bayesian Inference

- To a Bayesian, “inference” is then the process of updating prior expectations after seeing the data
- After we see the data: test scores


—  $p(\mu, \sigma^2 | \text{data}) = \frac{p(\text{data} | \mu, \sigma^2) \times p(\mu, \sigma^2)}{p(\text{data})}$




- This is equation follows directly from “Bayes’ Theorem”
- Note: this is also a distribution
- Known as a “posterior” distribution

# Pictorially

- You have a prior distribution on  $\sigma^2$  and  $\mu$ :




$\sigma^2$   
(variance of test scores)



$\mu$   
(average of test scores)


# Pictorially

- You see some test scores




$\sigma^2$

(variance of



$\mu$


range of test scores)



seem to average in the high 70's


# Pictorially

- Then use Bayes' to get a posterior distribution on  $\sigma^2$  and  $\mu$ :




$\sigma^2$

(variance of



$\mu$

(average of test scores)



much narrower  
after seeing data!

# That's the Bayesian Approach

- Come up with a generative process
- Which includes prior distributions on the quantities like to est
  - In our example, the variance and especially the mean
- See some data
- Use Bayes' Theorem and data to “update” the priors
  - This gives you a posterior dist
  - The posterior contains your estimate

# OK, So What Does This Have To Do W Text?

- Can apply same methodology to learning topics in text
- This is exactly what “LDA” does
  - Stands for “Latent Dirichlet Allocation”... fancy!
- First, we need a generative process
  - LDA will generate  $n$  random documents given a dictionary
  - Dictionary is of size num\_words
  - Best shown thru an example
  - In our example: dictionary will have: (0, “bad”) (1, “I”) (2, “can’t”) (3, “stand”) (4, “comp 215”) (5, “to”) (6, “leave”) (7, “love”) (8, “beer”) (9, “humanities”) (10, “classes”)



# LDA Step One

- Generate each of the  $k$  “topics”
  - Each topic is represented by a vector of probabilities
  - The  $w$ th entry in the vector is associated with the  $w$ th word in the dictionary
  - $\text{word\_probs}_t[w]$  is the probability that topic  $t$  would produce word  $w$
  - Vector is sampled from a Dirichlet (alpha) distribution
  - So, for each  $t$  in  $\{0 \dots k - 1\}$ ,  $\text{word\_probs}_t \sim \text{Dirichlet}(\alpha)$

# LDA Step One

- Generate each of the  $k$  “topics”
  - Each topic is represented by a vector of probabilities
  - The  $w$ th entry in the vector is associated with the  $w$ th word in the dictionary
  - $\text{word\_probs}_t[w]$  is the probability that topic  $t$  would produce word  $w$
  - Vector is sampled from a Dirichlet (alpha) distribution
  - So, for each  $t$  in  $\{0 \dots k - 1\}$ ,  $\text{word\_probs}_t \sim \text{Dirichlet}(\alpha)$
- Ex:  $k = 3$ 
  - $\text{word\_probs}_0 = (.2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0)$
  - $\text{word\_probs}_1 = (0, .2, .2, .2, 0, 0, 0, 0, 0, .2, .2)$
  - $\text{word\_probs}_2 = (0, .2, .2, 0, .2, 0, .2, .2, 0, 0, 0)$

## LDA Step Two

- Generate the topic proportions for each document
  - Each topic “controls” a subset of the words in a document
  - $\text{topic\_probs}_d[t]$  is the probability that an arbitrary word in document  $d$  will be controlled by topic  $t$
  - Vector is sampled from a Dirichlet (beta) distribution
  - So, for each  $d$  in  $\{0 \dots n - 1\}$ ,  $\text{topic\_probs}_d \sim \text{Dirichlet}(\text{beta})$

## LDA Step Two

- Generate the topic proportions for each document
  - Each topic “controls” a subset of the words in a document
  - $\text{topic\_probs}_d[t]$  is the probability that an arbitrary word in document  $d$  will be controlled by topic  $t$
  - Vector is sampled from a Dirichlet (beta) distribution
  - So, for each  $d$  in  $\{0 \dots n - 1\}$ ,  $\text{topic\_probs}_d \sim \text{Dirichlet}(\text{beta})$
- Ex:  $n = 4$ 
  - $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$
  - $\text{topic\_probs}_1 = (0.01, .98, 0.01)$
  - $\text{topic\_probs}_2 = (0.02, .49, .49)$
  - $\text{topic\_probs}_3 = (.98, 0.01, 0.01)$

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$ 
  - $t$  for word zero is...

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$ 
  - $t$  for word zero is zero, since we sampled (1, 0, 0) [there is a 1 in the zeroth entry]
  - So we generate the word using  $\text{word\_probs}_0 = (.2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0)$

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I”
  - $t$  for word zero is zero, since we sampled  $(1, 0, 0)$  [there is a 1 in the zeroth entry]
  - So we generate the word using  $\text{word\_probs}_0 = (.2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0)$
  - And we get  $(0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0)$ , which is equivalent to “I”



# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I”
  - Now onto the next word

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I”
  - $t$  for word one is zero, since we sampled  $(1, 0, 0)$  [there is a 1 in the zeroth entry]

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t”
  - $t$  for word one is zero, since we sampled  $(1, 0, 0)$  [there is a 1 in the zeroth entry]
  - So we generate the word using  $\text{word\_probs}_0 = (.2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0)$
  - And we get  $(0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0)$ , which is equivalent to “can’t”

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t”
  - Now onto the next word

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t”
  - $t$  for word two is zero, since we sampled  $(1, 0, 0)$  [there is a 1 in the zeroth entry]

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t stand”
  - $t$  for word two is zero, since we sampled  $(1, 0, 0)$  [there is a 1 in the zeroth entry]
  - So we generate the word using  $\text{word\_probs}_0 = (.2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0)$
  - And we get  $(0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0)$ , which is equivalent to “stand”

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t stand”
  - Onto next word

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t stand”
  - $t$  for word three is zero, since we sampled  $(1, 0, 0)$  [there is a 1 in the zeroth entry]



# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t stand bad”
  - $t$  for word three is zero, since we sampled  $(1, 0, 0)$  [there is a 1 in the zeroth entry]
  - So we generate the word using  $\text{word\_probs}_0 = (.2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0)$
  - And we get  $(1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)$ , which is equivalent to “bad”

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t stand bad”
  - Onto the last word in the document

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t stand bad”
  - $t$  for word three is zero, since we sampled  $(1, 0, 0)$  [there is a 1 in the zeroth entry]

# LDA Step Three

- Generate the words in each document
  - Each topic “controls” a subset of the words in a document
  - $\text{words\_in\_doc}_d[w]$  is the number of occurrences of word  $w$  in document  $d$
  - To get this vector, generate the words one-at-a-time
  - For a given word in doc  $d$ :
    - (1) Figure out the topic  $t$  that controls it by sampling from a Multinomial ( $\text{topic\_probs}_d, 1$ ) distribution
    - (2) Generate the word by sampling from a Multinomial ( $\text{word\_probs}_t, 1$ ) distribution
- Ex: doc 0...  $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$  “I can’t stand bad beer”
  - $t$  for word three is zero, since we sampled  $(1, 0, 0)$  [there is a 1 in the zeroth entry]
  - So we generate the word using  $\text{word\_probs}_0 = (.2, .2, .2, .2, 0, 0, 0, 0, .2, 0, 0)$
  - And we get  $(0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0)$ , which is equivalent to “beer”

## In The End... For Doc 0...

- text is “I can’t stand bad beer” (equiv. to “1 2 3 0 8”)
- $\text{topic\_probs}_0 = (.98, 0.01, 0.01)$
- $\text{words\_in\_doc}_0 = (1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0)$ 
  - Why? Word 0 appears once, word 1 appears once, word 4 zero times, etc.
- $\text{produced}_0 = (1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0)$   
 $(0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)$   
 $(0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)$ 
  - Why? Topic 0 (associated with first line) produced 5 words  
Those words were  $(1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0)$
  - Topic 1, topic 2 produced no words
  - “produced” always a matrix with  $\text{num\_words}$  cols,  $k$  rows

Repeat For Each Doc in the Corpus!

## For Example, Let's Look At Doc 2...

- $\text{topic\_probs}_2 = (.02, 0.49, 0.49)$
- Imagine that when we generate doc 2, we get:
  - Word 0: produced by topic 2, is 1 or “I”
  - Word 1: produced by topic 2, is 7 or “love”
  - Word 2: produced by topic 2, is 8 or “beer”
  - Word 3: produced by topic 1, is 1 or “I”
  - Word 4: produced by topic 1, is 2 or “can’t”
  - Word 5: produced by topic 2, is 7 or “love”
  - Word 6: produced by topic 1, is 9 or “humanities”
  - Word 7: produced by topic 1, is 10 or “classes”
- $\text{words\_in\_doc}_2 = (0, 2, 1, 0, 0, 0, 0, 2, 1, 1, 1)$
- $\text{produced}_2 =$   
     $(0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0)$   
     $(0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1)$   
     $(0, 1, 0, 0, 0, 0, 0, 2, 1, 0, 0)$

# OK We've Got Our Generative Process

- Now, someone gives us an actual, real-life corpus
  - This means we have got a dictionary
  - And we have  $\text{words\_in\_doc}_d$  for all  $d$  in  $\{0 \dots n - 1\}$
- Our goal is to figure out a posterior distribution on
  - $\text{word\_probs}_t$  for all  $t$  in  $\{0 \dots k - 1\}$
  - $\text{topic\_probs}_d$  for all  $d$  in  $\{0 \dots n - 1\}$
  - $\text{produced}_d$  for all  $d$  in  $\{0 \dots n - 1\}$
- Why?
  - The hope is this will reveal something about the corpus
  - For example, what the different topics in the corpus are
  - And what topics are present in each document



# As In All Applications of Bayesian ML

- The posterior is derived using Bayes' Theorem...

$$\begin{aligned} p(\text{word\_probs}_{all}, \text{topic\_probs}_{all}, \text{produced}_{all} \mid \text{words\_in\_doc}_{all}) = \\ p(\text{words\_in\_doc}_{all} \mid \text{word\_probs}_{all}, \text{topic\_probs}_{all}, \text{produced}_{all}) \times \\ p(\text{word\_probs}_{all}, \text{topic\_probs}_{all}, \text{produced}_{all}) / p(\text{words\_in\_doc}_{all}) \end{aligned}$$

# We Can Write This Formula

- But what does this do for us?
- In the simple “guess the test score average” case...
  - We had a couple of posterior distributions that we could plot
  - It was quite intuitive
- Not the case here!
  - Got all kinds of weird, multi-dimensions distributions
  - How to proceed?

# With More Complicated Models Such as LDA

- It is common to resort to something called “MCMC”
  - stands for “Markov Chain Monte Carlo”
- MCMC is a class of algorithms
  - That can often be used to draw samples from even the most complex distributions
  - Many of the ideas behind MCMC came out the the Manhattan project
- Using MCMC, we can actually draw samples from
$$p(\text{word\_probs}_{all}, \text{topic\_probs}_{all}, \text{produced}_{all} \mid \text{words\_in\_doc}_{all})$$
- Each “sample” will contain everything in the model!
  - $\text{word\_probs}_t$  for all  $t$  in  $\{0 \dots k - 1\}$
  - $\text{topic\_probs}_d$  for all  $d$  in  $\{0 \dots n - 1\}$
  - $\text{produced}_d$  for all  $d$  in  $\{0 \dots n - 1\}$

# Why Useful?

- Can take one sample, use it as a representative set of values
- Or take many samples, use to get a summary of the distribution

# How Does MCMC Work?

- Could spend an entire semester on MCMC alone
- Many flavors of MCMC
- For LDA, we'll employ a simple MCMC methodology
  - A “Gibbs Sampler”

# In the Case of LDA

- Here is pseudo-code for the Gibbs sampler:

Choose consistent, initial values for  $\text{word\_probs}_{all}$ ,  $\text{topic\_probs}_{all}$ ,  $\text{produced}_{all}$

For  $i = 1$  to  $\text{num\_iters}$  do:

For all  $t$  in  $\{0 \dots k - 1\}$ :

$\text{word\_probs}_t \sim p(\text{word\_probs}_t \mid \text{topic\_probs}_{all}, \text{produced}_{all}, \text{words\_in\_doc}_{all})$

For all  $d$  in  $\{0 \dots n - 1\}$ :

$\text{topic\_probs}_d \sim p(\text{topic\_probs}_d \mid \text{word\_probs}_{all}, \text{produced}_{all}, \text{words\_in\_doc}_{all})$

For all  $d$  in  $\{0 \dots n - 1\}$ :

$\text{produced}_d \sim p(\text{produced}_d \mid \text{word\_probs}_{all}, \text{topic\_probs}_{all}, \text{words\_in\_doc}_{all})$

- Run this loop for awhile

— Then at any point, stop

— Current  $\text{word\_probs}_{all}$ ,  $\text{topic\_probs}_{all}$ ,  $\text{produced}_{all}$  are a sample from posterior!

## So All That's Left

- Is to give pseudo-code for each of the three steps

# Word\_Probs

- To sample each  $\text{word\_probs}_t$ 
  - Create a vector called “counter”, where counter is  $\sum_d \text{produced}_{d,t}$
  - Note that  $\text{produced}_{d,t}$  denotes the  $t$ th row of the produced matrix for doc  $d$
  - This counts the number of times topic  $t$  produced each word  $w$
  - Then  $\text{word\_probs}_t \sim \text{Dirichlet}(\text{counter} + \alpha)$ 

↑  
Constant used in generative process  
We assume we know this,  
or can guess it



# Word\_Probs

- To sample each  $\text{word\_probs}_t$ 
  - Create a vector called “counter”, where counter is  $\sum_d \text{produced}_{d,t}$
  - Note that  $\text{produced}_{d,t}$  denotes the  $t$ th row of the produced matrix for doc  $d$
  - This counts the number of times topic  $t$  produced each word  $w$
  - Then  $\text{word\_probs}_t \sim \text{Dirichlet}(\text{counter} + \text{alpha})$
- Intuitively
  - You are “guessing” the probability topic  $t$  will produce each word
  - If  $\text{counter}[w]$  is large, then topic  $t$  produced  $w$  quite often in practice...
  - And the Dirichlet is then likely to give that word a high probability

# Topic\_Probs

- To sample each  $\text{topic\_probs}_d$ 
  - Create a vector “counter”, where  $\text{counter}[t]$  is  $\sum_w \text{produced}_{d,t}[w]$
  - That is,  $\text{counter}[t]$  is the sum over the  $t$ th row in the  $\text{produced}_d$  matrix
  - This counts the number of times topic  $t$  was used to produce a word in  $d$
  - Then  $\text{topic\_probs}_d \sim \text{Dirichlet}(\text{counter} + \text{beta})$

↑  
Again, constant used in generative process  
We assume we know this,  
or can guess it

# Topic\_Probs

- To sample each  $\text{topic\_probs}_d$ 
  - Create a vector “counter”, where  $\text{counter}[t]$  is  $\sum_w \text{produced}_{d,t}[w]$
  - That is,  $\text{counter}[t]$  is the sum over the  $t$ th row in the  $\text{produced}_d$  matrix
  - This counts the number of times topic  $t$  was used to produce a word in  $d$
  - Then  $\text{topic\_probs}_d \sim \text{Dirichlet}(\text{counter} + \text{beta})$
- Intuitively
  - You are “guessing” the probability a word in  $d$  will come from each topic
  - If  $\text{counter}[t]$  is large, then topic  $t$  produced a lot of words in  $d$ ...
  - And the Dirichlet is then likely to give that topic a high probability

# Produced

- To sample each  $\text{produced}_d$ 
  - For each word  $w$ :
    - (1) Create a vector “probs”, where  $\text{probs}[t] = \text{word\_probs}_t[w] \times \text{topic\_probs}_d[t]$
    - (2) Normalize probs
    - (3) Then ( $w$ th column in  $\text{produced}_d$ )  $\sim \text{Multinomial}(\text{words\_in\_doc}_d, \text{probs})$

# Produced

- To sample each produced<sub>*d*</sub>
  - For each word *w*:
    - (1) Create a vector “probs”, where  $\text{probs}[t] = \text{word\_probs}_t[w] \times \text{topic\_probs}_d[t]$
    - (2) Normalize probs
    - (3) Then (*w*th column in produced<sub>*d*</sub>)  $\sim \text{Multinomial}(\text{words\_in\_doc}_d, \text{probs})$
- Intuitively
  - “probs” is telling you how likely each topic is to be the one responsible for an occurrence of *w* in *d*
  - Then you are using the multinomial to guess how many occurs of *w* each topic is responsible for
  - $\text{probs}[t]$  large means *t* expectedly responsible for more occurs of *w*

This Completes LDA. Questions?